

## HIVE INSTALLATION:

```
hdp123@stud-OptiPlex-3060: /usr/local/hive/bin
stud@stud-OptiPlex-3060:~$ echo $JAVA_HOME
/usr/lib/jvm/java-1.8.0-openjdk-amd64
stud@stud-OptiPlex-3060:~$ su - hdp123
Password:
hdp123@stud-OptiPlex-3060:~$ start-dfs.sh
WARNING: An illegal reflective access operation has occurred
WARNING: Illegal reflective access by org.apache.hadoop.security.authentication.util.KerberosUtil (file:/usr/local/hadoop/share/hadoop/common/lib/hadoop-auth-2.10.1.jar) to method sun.security.krb5.Config.getInstance()
WARNING: Please consider reporting this to the maintainers of org.apache.hadoop.security.authentication.util.KerberosUtil
WARNING: Use --illegal-access=warn to enable warnings of further illegal reflective access operations
WARNING: All illegal access operations will be denied in a future release
23/04/27 16:48:06 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Starting namenodes on [localhost]
localhost: starting namenode, logging to /usr/local/hadoop/logs/hadoop-hdp123-namenode-stud-OptiPlex-3060.out
localhost: WARNING: An illegal reflective access operation has occurred
localhost: WARNING: Illegal reflective access by org.apache.hadoop.security.authentication.util.KerberosUtil (file:/usr/local/hadoop/share/hadoop/common/lib/hadoop-auth-2.10.1.jar) to method sun.security.krb5.Config.getInstance()
localhost: WARNING: Please consider reporting this to the maintainers of org.apache.hadoop.security.authentication.util.KerberosUtil
localhost: WARNING: Use --illegal-access=warn to enable warnings of further illegal reflective access operations
localhost: WARNING: All illegal access operations will be denied in a future release
localhost: starting datanode, logging to /usr/local/hadoop/logs/hadoop-hdp123-datanode-stud-OptiPlex-3060.out
```

```
hdp123@stud-OptiPlex-3060: /usr/local/hive/bin
hdp123@stud-OptiPlex-3060:~$ jps
8480 NameNode
9062 ResourceManager
9195 NodeManager
8683 DataNode
8891 SecondaryNameNode
10075 Jps
hdp123@stud-OptiPlex-3060:~$ cd /usr/local/hadoop
hdp123@stud-OptiPlex-3060: /usr/local/hadoop$ ls
bin  include  libexec  logs  README.txt  share
etc  lib      LICENSE.txt  NOTICE.txt  sbin
hdp123@stud-OptiPlex-3060: /usr/local/hadoop$ cd bin
hdp123@stud-OptiPlex-3060: /usr/local/hadoop/bin$ hdfs dfs -mkdir -p /usr/hive/warehouse
23/04/27 16:58:59 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
hdp123@stud-OptiPlex-3060: /usr/local/hadoop/bin$ hdfs dfs -mkdir -p /tmp/hive
23/04/27 16:59:31 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
hdp123@stud-OptiPlex-3060: /usr/local/hadoop/bin$ hdfs dfs -chmod 777 /tmp
23/04/27 17:00:16 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
hdp123@stud-OptiPlex-3060: /usr/local/hadoop/bin$ hdfs dfs -chmod 777 /usr/hive/warehouse
23/04/27 17:00:34 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
```

```
hdp123@stud-OptiPlex-3060: /usr/local/hive/bin
hdp123@stud-OptiPlex-3060: /usr/local/hadoop/bin$ hdfs dfs -chmod 777 /tmp/hive
23/04/27 17:01:04 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
chmod: '/tmp/hive': No such file or directory
hdp123@stud-OptiPlex-3060: /usr/local/hadoop/bin$ hdfs dfs -chmod 777 /tmp/hive
23/04/27 17:01:10 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
hdp123@stud-OptiPlex-3060: /usr/local/hadoop/bin$ cd ..
hdp123@stud-OptiPlex-3060: /usr/local/hadoop$ cd ..
hdp123@stud-OptiPlex-3060: /usr/local$ cd hive
hdp123@stud-OptiPlex-3060: /usr/local/hive$ ls
bin  binary-package-licenses  conf  examples  hcatlog  jdbc  lib  LICENSE  NOTICE  RELEASE_NOTES.txt  scripts
hdp123@stud-OptiPlex-3060: /usr/local/hive$ cd bin
hdp123@stud-OptiPlex-3060: /usr/local/hive/bin$ ls
beeline  ext  hive  hive-conf  hiveserver2  hplsql  metatool  schematool
hdp123@stud-OptiPlex-3060: /usr/local/hive/bin$ SHIVE_HOME/bin/schematool -dbType derby -initSchema
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/usr/local/hive/lib/log4j-slf4j-impl-2.6.2.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/usr/local/hadoop/share/hadoop/common/lib/slf4j-log4j12-1.7.25.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]
```

```
hdp123@stud-OptiPlex-3060: /usr/local/hive/bin
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]
Metastore connection URL: jdbc:derby:::databaseName=metastore_db;create=true
Metastore Connection Driver : org.apache.derby.jdbc.EmbeddedDriver
Metastore connection User: APP
Starting metastore schema initialization to 2.3.0
Initialization script hive-schema-2.3.0.derby.sql
Initialization script completed
schemaTool completed
hdp123@stud-OptiPlex-3060: /usr/local/hive/bin$ ls
beeline derby.log ext hive hive-config.sh hiveserver2 hplsql metastore_db metatool schenatool
hdp123@stud-OptiPlex-3060: /usr/local/hive/bin$ ./hive
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/usr/local/hive/lib/log4j-slf4j-impl-2.6.2.jar!/org.slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/usr/local/hadoop/share/hadoop/common/lib/slf4j-log4j12-1.7.25.jar!/org.slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]

Logging initialized using configuration in jar:file:/usr/local/hive/lib/hive-common-2.3.9.jar!/hive-log4j2.properties Async: true
Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.
hive> show databases;
```

## Adding path of hive in bashrc:

```
hdp123@stud-OptiPlex-3060: /usr/local
GNU nano 6.2 /home/hdp123/.bashrc
# sources /etc/bash.bashrc).
if ! shopt -oq posix; then
  if [ -f /usr/share/bash-completion/bash_completion ]; then
    . /usr/share/bash-completion/bash_completion
  elif [ -f /etc/bash_completion ]; then
    . /etc/bash_completion
  fi
fi
export JAVA_HOME=/usr/lib/jvm/java-11-openjdk-amd64
export HADOOP_HOME=/usr/local/hadoop
export PATH=$PATH:$HADOOP_HOME/bin
export PATH=$PATH:$HADOOP_HOME/sbin
export HADOOP_MAPRED_HOME=$HADOOP_HOME
export HADOOP_COMMON_HOME=$HADOOP_HOME
export HADOOP_HDFS_HOME=$HADOOP_HOME
export YARN_HOME=$HADOOP_HOME
export HADOOP_COMMON_LIB_NATIVE_DIR=$HADOOP_HOME/lib/native
export HADOOP_OPTS="-Djava.library.path=$HADOOP_HOME/lib"

export HIVE_HOME=/usr/local/hive
|
^G Help      ^O Write Out ^W Where Is  ^K Cut       ^T Execute  ^C Location  M-U Undo     M-A Set Mark
^X Exit      ^R Read File ^_ Replace   ^U Paste     ^J Justify  ^/_ Go To Line M-E Redo     M-6 Copy
```

## Hive Terminal:

```
hdp123@stud-OptiPlex-3060: /usr/local/hive/bin
hive> show tables;
OK
emp
Time taken: 5.405 seconds, Fetched: 1 row(s)
hive> drop table emp;
OK
Time taken: 0.399 seconds
hive> create table emp(ename string,esal int) row format delimited fields terminated by ',' stored as textfile;
OK
Time taken: 0.487 seconds
hive> load data local inpath '/home/stud/Desktop/emp_details.txt' into table emp;
Loading data to table default.emp
OK
Time taken: 0.365 seconds
hive> select * from emp;
OK
sachin 10000
ayush 50000
yuvraj 70000
yogesh 90000
rahul 11000
mark 20000
stocks 40000
Time taken: 0.74 seconds, Fetched: 7 row(s)
hive>
```

## Creating an external hive table:

```
hdp123@stud-OptiPlex-3060:/usr/local$ hdfs dfs -mkdir /user/hive/flight
23/05/02 14:31:31 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
hdp123@stud-OptiPlex-3060:/usr/local$ hdfs dfs -put /home/stud/Downloads/flight_ext.csv /user/hive/flight
23/05/02 14:34:41 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
hdp123@stud-OptiPlex-3060:/usr/local$ hdfs dfs -ls /user/hive/flight
23/05/02 14:34:58 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 1 items
-rw-r--r-- 1 hdp123 supergroup 3842 2023-05-02 14:34 /user/hive/flight/flight_ext.csv
```

```
hdp123@stud-OptiPlex-3060:/usr/local/hive/bin$ ./hive
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/usr/local/hive/lib/log4j-slf4j-impl-2.6.2.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/usr/local/hadoop/share/hadoop/common/lib/slf4j-log4j12-1.7.25.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]

Logging initialized using configuration in jar:file:/usr/local/hive/lib/hive-common-2.3.9.jar!/hive-log4j2.properties Async: true
Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.
hive> create external table flight_ext(Year int, Month int, DayOfMonth int, DayOfWeek int, DepTime int, CRSDepTime int, ArrTime int, CRSArrTime int, UniqueCarrier string, FlightNum int, TailNum string, ActualElapsedTime int, CRSElapsedTime int, AirTime int, ArrDelay int, DepDelay int, Origin string, Dest string, Distance int, TaxiIn int, TaxiOut int, Cancelled int, CancellationCode string, Diverted string, CarrierDelay int, WeatherDelay int, NASDelay int, SecurityDelay int, LateAircraftDelay int) row format delimited fields terminated by ',' location '/user/hive/flight';
OK
Time taken: 20.014 seconds
hive> show tables;
OK
emp
flight
flight_ext
flight_fare
Time taken: 0.078 seconds, Fetched: 4 row(s)
```

To find the details of table in hive:

```
hive> describe formatted flight_ext;
OK
# col_name          data_type          comment
year                int
month               int
dayofmonth           int
dayofweek            int
deptime             int
crsdeptime           int
arrtime             int
crsarrrtime         int
uniquecarrier       string
flightnum           int
tailnum             string
actualelapsedtime   int
crselapsedtime      int
airtime             int
arrdelay            int
depdelay            int
origin              string
dest                string
distance            int
taxiin              int
taxiout             int
cancelled           int
cancellationcode    string
diverted            string
carrierdelay        int
weatherdelay        int
nasdelay            int
securitydelay       int
lateaircraftdelay   int
```

Count the number of rows in table:

```
hive> select count(*) from flight_ext;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e.
spark, tez) or using Hive 1.X releases.
Query ID = hdp123_20230502143924_c7d1e875-d69a-4964-8cba-18477918dc69
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1683014056661_0002, Tracking URL = http://stud-OptiPlex-3060:8088/proxy/application_1683014056661_0002/
Kill Command = /usr/local/hadoop/bin/hadoop job -kill job_1683014056661_0002
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2023-05-02 14:39:30,423 Stage-1 map = 0%, reduce = 0%
2023-05-02 14:39:34,513 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 1.36 sec
2023-05-02 14:39:39,592 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 2.79 sec
MapReduce Total cumulative CPU time: 2 seconds 790 msec
Ended Job = job_1683014056661_0002
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 2.79 sec HDFS Read: 15222 HDFS Write: 102 SUCCESS
Total MapReduce CPU Time Spent: 2 seconds 790 msec
OK
39
Time taken: 15.681 seconds, Fetched: 1 row(s)
hive>
```

## Creating an internal table in hive:

### Creating flight table:

```
hdp123@stud-OptiPlex-3060: /usr/local/hive/bin
Time taken: 3.076 seconds, Fetched: 2 row(s)
hive> show tables;
OK
emp
Time taken: 0.043 seconds, Fetched: 1 row(s)
hive> create table flight(Year int, Month int, DayOfMonth int, DayOfWeek int, DepTime int,
CRSDepTime int, ArrTime int, CRSArrTime int, UniqueCarrier string, FlightNum int, TailNum
string, ActualElapsedTime int, CRSElapsedTime int, AirTime int, ArrDelay int, DepDelay int,
Origin string, Dest string, Distance int, TaxiIn int, TaxiOut int, Cancelled int, Cancellat
ionCode string, Diverted string, CarrierDelay int, WeatherDelay int, NASDelay int, Securi
tyDelay int, LateAircraftDelay int) row format delimited fields terminated by ',';
OK
Time taken: 0.713 seconds
hive> load data local inpath '/home/stud/Downloads/flight_data.csv' into table flight;
Loading data to table default.flight
OK
Time taken: 0.979 seconds
hive> show tables;
OK
emp
flight
Time taken: 0.022 seconds, Fetched: 2 row(s)
hive> select * from flight;
OK
```

### Creating flight\_fare table:

```
hdp123@stud-OptiPlex-3060: /usr/local/hive/bin
Time taken: 0.021 seconds, Fetched: 2 row(s)
hive> create table flight_fare( FlightNum int, Fare int, Distance int) row format delimited fields terminated by ',';
OK
Time taken: 0.074 seconds
hive> show tables;
OK
emp
flight
flight_fare
Time taken: 0.02 seconds, Fetched: 3 row(s)
hive> load data local inpath '/home/stud/Downloads/flight_fare_new.csv' into table flight_fare;
Loading data to table default.flight_fare
OK
Time taken: 0.352 seconds
```

```
hive> select FlightNum, Origin, Dest from flight limit 6;
OK
324      SEA      SJC
572      SEA      PSP
511      SAN      SEA
376      SEA      GEG
729      TUS      SEA
283      LAX      SEA
Time taken: 0.108 seconds, Fetched: 6 row(s)
hive> select FlightNum, Fare, Distance from flight_fare limit 6;
OK
324      5181      697
572      7905      987
511      8781      1050
376      7852      224
729      9009      1216
283      6250      954
Time taken: 0.094 seconds, Fetched: 6 row(s)
hive> █
```



## Join operation on flight and flight\_fare table:

```
Activities Terminal May 2 14:24 hdp123@stud-OptiPlex-3060: /usr/local/hive/bin

hive> select flight.Origin, flight.Dest, flight_fare.fare, flight_fare.Distance from flight join flight_fare on flight.FlightNum=flight_fare.FlightNum;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.
Query ID = hdp123_20230502141316_5a35aed2-d5d5-4c9c-a5fe-681ec37204ea
Total jobs = 1
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/usr/local/hive/lib/log4j-slf4j-impl-2.6.2.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/usr/local/hadoop/share/hadoop/common/lib/slf4j-log4j12-1.7.25.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]
2023-05-02 14:13:24 Starting to launch local task to process map join; maximum memory = 477626368
2023-05-02 14:13:25 Dump the side-table for tag: 1 with group count: 47 into file: file:/tmp/hdp123/639da944-04a1-4db5-9ced-63b7468f8b13/hive_2023-05-02_14-13-16_922_5956281314495289346-1/-local-10004/HashTable-Stage-3/MapJoin-mapfile01--.hashtable
2023-05-02 14:13:25 Uploaded 1 File to: file:/tmp/hdp123/639da944-04a1-4db5-9ced-63b7468f8b13/hive_2023-05-02_14-13-16_922_5956281314495289346-1/-local-10004/HashTable-Stage-3/MapJoin-mapfile01--.hashtable (1522 bytes)
2023-05-02 14:13:25 End of local task; Time Taken: 0.96 sec.
Execution completed successfully
MapredLocal task succeeded
Launching Job 1 out of 1
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = job_1683014056661_0001, Tracking URL = http://stud-OptiPlex-3060:8088/proxy/application_1683014056661_0001/
Kill Command = /usr/local/hadoop/bin/hadoop job -kill job_1683014056661_0001
Hadoop job information for Stage-3: number of mappers: 1; number of reducers: 0
2023-05-02 14:13:33,983 Stage-3 map = 0%, reduce = 0%
2023-05-02 14:13:38,085 Stage-3 map = 100%, reduce = 0%, Cumulative CPU 1.96 sec
MapReduce Total cumulative CPU time: 1 seconds 960 msec
Ended Job = job_1683014056661_0001
MapReduce Jobs Launched:
Stage-Stage-3: Map: 1 Cumulative CPU: 1.96 sec HDFS Read: 13263 HDFS Write: 1550 SUCCESS
Total MapReduce CPU Time Spent: 1 seconds 960 msec
OK
SEA SJG 5181 697
SEA PSP 7905 987
SAN SEA 8781 1050
SEA GEG 7852 224
TUS SEA 9009 1216
LAX SEA 6250 954
```

## To find average departure delay per day in 2008:

```
hive> select avg(depdelay) from flight;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.
Query ID = hdp123_20230502152931_7721f663-a16b-4624-8b5f-63ad67c70ddf
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1683021424703_0001, Tracking URL = http://stud-OptiPlex-3060:8088/proxy/application_1683021424703_0001/
Kill Command = /usr/local/hadoop/bin/hadoop job -kill job_1683021424703_0001
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2023-05-02 15:29:39,834 Stage-1 map = 0%, reduce = 0%
2023-05-02 15:29:43,943 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 1.15 sec
2023-05-02 15:29:49,020 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 2.63 sec
MapReduce Total cumulative CPU time: 2 seconds 630 msec
Ended Job = job_1683021424703_0001
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 2.63 sec HDFS Read: 16632 HDFS Write: 118 SUCCESS
Total MapReduce CPU Time Spent: 2 seconds 630 msec
OK
10.021739130434783
Time taken: 20.024 seconds, Fetched: 1 row(s)
hive>
```